



## LECTURE 3

## FUNDAMENTAL GPU ALGORITHMS

- REDUCE
- SCAN
- HISTOGRAM



## DIGGING HOLES AGAIN



## QUIZ


# STEPS?

TOTAL  
AMOUNT  
OF work?




## REDUCE

$$1 + 2 + 3 + 4 + \dots$$




## REDUCE

$$1 + 2 + 3 + 4 + \dots$$



## REDUCE: INPUTS

- 1) SET OF ELEMENTS
- 2) REDUCTION OPERATOR




## SERIAL IMPLEMENTATION OF REDUCE

### SERIAL CODE

```
sum = 0
for (i=0; i < elts.len(); i++) {
    sum = sum + elts[i]
}
return sum
```

### REDUCE







### Quiz

Which are true about a serial reduce code running on an input of size  $n$ ?

- ☐ It takes  $n$  operations.
- ☐ It takes  $n-1$  operations.
- ☐ Its work complexity is  $O(n)$
- ☐ Its step complexity is  $O(1)$



### PARALLEL REDUCE




### Quiz

How do you rewrite  $(a+b)+c+d$  to allow parallel execution?



(Use parens to show grouping.)



## STEP COMPLEXITY OF PARALLEL REDUCTION

N	STEPS
2	1
4	2
8	3
...	...



## STEP COMPLEXITY OF PARALLEL REDUCTION

N	STEPS
2	1
4	2
8	3
...	...

Quiz

☐  $\sqrt{n}$

☐  $\log_2 n$

☐  $n$


☐  $n \log_2 n$



## REDUCING 1M ELEMENTS

① 1024 BLOCKS x 1024 THREADS

② 1 BLOCK x 1024 THREADS



```

global__ void global_reduce_kernel(float * d_out, float * d_in)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int tid = threadIdx.x;

    // do reduction in global mem
    for (unsigned int s = blockDim.x / 2; s > 0; s >>= 1)
    {
        if (tid < s)
        {
            d_in[myId] += d_in[myId + s];
        }
        __syncthreads(); // make sure all adds at one stage are done!
    }

    // only thread 0 writes result for this block back to global mem
    if (tid == 0)
    {
        d_out[blockIdx.x] = d_in[myId];
    }
}
    
```

U:--- reduce.cu 2% L16 (C++/1 A

```

__global__ void shmem_reduce_kernel(float * d_out, const float * d_in)
{
    // sdata is allocated in the kernel call: 3rd arg to <<<b, t, shmem>>>
    extern __shared__ float sdata[];

    int myId = threadIdx.x + blockDim.x * blockDim.x;
    int tid = threadIdx.x;

    // load shared mem from global mem
    sdata[tid] = d_in[myId];
    __syncthreads(); // make sure entire block is loaded!

    // do reduction in shared mem
    for (unsigned int s = blockDim.x / 2; s > 0; s >>= 1)
    {
        if (tid < s)
        {
            sdata[tid] += sdata[tid + s];
        }
        __syncthreads(); // make sure all adds at one stage are done!
    }
}

```

U:--- reduce.cu 13% L38 (C++/1 Abbrev)

```

int tid = threadIdx.x;

// load shared mem from global mem
sdata[tid] = d_in[myId];
__syncthreads(); // make sure entire block is loaded!

// do reduction in shared mem
for (unsigned int s = blockDim.x / 2; s > 0; s >>= 1)
{
    if (tid < s)
    {
        sdata[tid] += sdata[tid + s];
    }
    __syncthreads(); // make sure all adds at one stage are done!
}

// only thread 0 writes result for this block back to global mem
if (tid == 0)
{
    d_out[blockIdx.x] = sdata[0];
}
}

```

U:--- reduce.cu 17% L44 (C++/1 Abbrev)

## SHARED VS GLOBAL MEMORY BANDWIDTH

THE GLOBAL MEMORY VERSION USES



TIMES AS MUCH GLOBAL MEM BW AS  
THE SHARED MEM VERSION?

## SCAN

— EXAMPLE

INPUT: 1 2 3 4

OPERATION: ADD

OUTPUT: 1 3 6 10

## SCAN


### — EXAMPLE

INPUT: 1 2 3 4

OPERATION: ADD

OUTPUT: 1 3 6 10

- ADDRESSES SET OF PROBLEMS OTHERWISE DIFFICULT TO PARALLELIZE
- NOT USEFUL IN SERIAL WORLD BUT VERY USEFUL IN PARALLEL
- TODAY: EXPLAINING WHAT + HOW  
BUT NOT WHY (NEXT LECTURE)



TRANSACTION	BALANCE
\$ 20	20
5	25
- 11	14
- 9	5
- 3	2
15	17
INPUT	OUTPUT

### INPUTS TO SCAN

- INPUT ARRAY
- BINARY ASSOCIATIVE OPERATOR } LIKE REDUCE
- IDENTITY ELEMENT [  $I \text{ op } a = a$  ]

OP	I	BECAUSE
+	0	$0 + a = a$
min (on unsigned chars)	0xFF	$\text{min}(0xFF, a) = a$

### QUIZ

WHAT IS THE IDENTITY FOR ...

Multiply

Logical or

Logical and

What Scan Does

Input: array A, operator  $\oplus$ , identity I

$[a_0 \ a_1 \ a_2 \ a_3 \ \dots \ a_{n-1}]$  INPUT

$[I \ a_0 \ a_0 \oplus a_1 \ a_0 \oplus a_1 \oplus a_2 \ \dots \ a_0 \oplus a_1 \oplus a_2 \oplus \dots \oplus a_{n-2}]$  OUTPUT

What Scan Does

Input: array A, operator  $\oplus$ , identity I  
PLUS  $\emptyset$

$[a_0 \ a_1 \ a_2 \ a_3 \ \dots \ a_{n-1}]$  INPUT

$[I \ a_0 \ a_0 \oplus a_1 \ a_0 \oplus a_1 \oplus a_2 \ \dots \ a_0 \oplus a_1 \oplus a_2 \oplus \dots \oplus a_{n-2}]$  OUTPUT

IN  $[3 \ 1 \ 4 \ 1 \ 5 \ 9]$

OUT  $[0 \ 3 \ 4 \ 8 \ 9 \ 14]$

QUIZ

MAX-SCAN  
ON UNSIGNED  
INTS

$[3 \ 1 \ 4 \ 1 \ 5 \ 9]$

OUTPUT?  $[\square \ \square \ \square \ \square \ \square \ \square]$

SERIAL IMPLEMENTATION OF SCAN

```
int acc = identity;
for (i=0; i < elements.length(); i++) {
    acc = acc op element[i];
    out[i] = acc;
}
```

## SERIAL IMPLEMENTATION OF SCAN

INCLUSIVE

```
int acc = identity;
for (i=0; i < elements.length(); i++) {
    acc = acc op element[i];
    out[i] = acc;
}
```

QUIZ: CONVERT  
TO EXCLUSIVE  
SCAN.

## INCLUSIVE VS EXCLUSIVE SCAN

INPUT: [13 7 16 21 8 20 13 12]

EXCLUSIVE  
SCAN  
OUTPUT

INCLUSIVE  
SCAN  
OUTPUT

## INCLUSIVE VS EXCLUSIVE SCAN

INPUT: [13 7 16 21 8 20 13 12]

EXCLUSIVE  
SCAN  
OUTPUT: [0 13 20 36 57 65 85 98] OUTPUT: ALL  
ELEMENTS  
BEFORE, NOT  
CURRENT ELT.

INCLUSIVE  
SCAN  
OUTPUT: [13 20 36 57 65 85 98 110] OUTPUT: ALL  
ELEMENTS  
BEFORE AND  
CURRENT ELT.

## SERIAL IMPLEMENTATION OF SCAN

INCLUSIVE

```
int acc = identity;
for (i=0; i < elements.length(); i++) {
    acc = acc op element[i];
    out[i] = acc;
}
```

WORK? n

STEPS? n



WHY SCAN IS USEFUL FOR PARALLELIZATION

INPUTS      0   0   0   0   0  
 OUTPUTS    0   0   0   0   0

WHY SCAN IS USEFUL FOR PARALLELIZATION

INPUTS      0   0   0   0   0  
 OUTPUTS    0   0   0   0   0

INCLUSIVE SCAN EXAMPLE, REVISITED

IN: [3 1 4 1 5 9]

OUT: [3 4 8 9 14 23]

INCLUSIVE SCAN EXAMPLE, REVISITED

IN: [3 1 4 1 5 9]

OUT: [3 4 8 9 14 23]

TWO PARALLEL SCAN ALGORITHMS

	MORE STEP-EFFICIENT	MORE WORK-EFFICIENT
HILLIS + STEELE	X	
BLELLOCH		X

INCLUSIVE SCAN EXAMPLE, REVISITED

IN: [3 1 4 1 5 9]

OUT: [3 4 8 9 14 23]


QUIZ

	CONSTANT $O(1)$	$O(\log n)$	LINEAR $O(n)$	$O(n^2)$
STEPS?	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
WORK?	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>


# HILLIS/STEELE INCLUSIVE SCAN



# HILLIS/STEELE INCLUSIVE SCAN



# HILLIS/STEELE INCLUSIVE SCAN




STARTING WITH STEP 0:  
ON STEP  $i$ , ADD YOURSELF  
TO YOUR  $2^i$  LEFT NEIGHBOR

# BRELOCH SCAN - REDUCE/DOWNSWEEP - EXCLUSIVE




# BRELOCH SCAN - REDUCE/DOWNSWEEP - EXCLUSIVE





## Q112

MAX SCAN  
USING  
REDUCE/  
DOWNSWEEP



# BRELOCH SCAN - REDUCE/DOWNSWEEP - EXCLUSIVE



## QUIZ

SERIAL

HILLIS  
STEELE

BLELOCH

512 ELT. VECTOR  
512 PROCESSORS



1M ELT. VECTOR  
512 PROCESSORS




128 K ELT. VECTOR  
1 PROCESSOR



## HISTOGRAM

160 cm  
175 cm  
152 cm  
...



## HISTOGRAM

COUNT


IN: HISTOGRAM  
OUT: CDF

OPERATION



160 cm  
175 cm  
152 cm

CUMULATIVE  
DISTRIBUTION  
FUNCTION



## HISTOGRAM

COUNT


IN: HISTOGRAM  
OUT: CDF

OPERATION



160 cm  
175 cm  
152 cm

CUMULATIVE  
DISTRIBUTION  
FUNCTION



# SERIAL ALGORITHM: HISTOGRAM

```
for (i=0; i<BIN_COUNT; i++)
```

```
    result[i] = 0;
```

```
for (i=0; i<BIN_COUNT; i++)
```

```
    result[computeBin(measurements[i])]++;
```

↳ TO WHICH BIN DOES THIS MEASUREMENT BELONG?

INPUT:

155

150

175

170

0	<150
0	150-165
0	165-180
0	>180

# SERIAL ALGORITHM: HISTOGRAM

```
for (i=0; i<BIN_COUNT; i++)
```

```
    result[i] = 0;
```

```
for (i=0; i<BIN_COUNT; i++)
```

```
    result[computeBin(measurements[i])]++;
```

↳ TO WHICH BIN DOES THIS MEASUREMENT BELONG?

INPUT:

155

150

175

170

0	<150
2	150-165
2	165-180
0	>180

QUIZ

n measurements

b bins

MAXIMUM # OF MEASUREMENTS/BIN

AVERAGE # OF MEASUREMENTS/BIN

# SERIAL ALGORITHM: HISTOGRAM

```
for (i=0; i<BIN_COUNT; i++)
```

```
    result[i] = 0;
```

```
for (i=0; i<measurements.size(); i++)
```

```
    result[computeBin(measurements[i])]++;
```

```


{
    int r = 0;
    for (int i = 0; i < bits; i++)
    {
        int bit = (w & (1 << i)) >> i;
        r |= bit << (bits - i - 1);
    }
    return r;
}

__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}


__global__ void simple_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
}

--:-- histo.cu 6% 129 (C++/1 Abbrev)
    
```


WHY THE OBVIOUS METHOD DOESN'T WORK



WHY THE OBVIOUS METHOD DOESN'T WORK



METHOD 1: ACCUMULATE USING ATOMICS



```
}
return r;
}

__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockDim.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}

__global__ void simple_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockDim.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    atomicAdd(&d_bins[myBin], 1);
}

int main(int argc, char **argv)
{
    histo.cu 10% L37 (C++/1 Abbrev)
}
```

## QUIZ

- Histogram with 1M elements
- You can choose # of bins




PER-THREAD PRIVATIZED (LOCAL) HISTOGRAMS, THEN REDUCE  
128 ITEMS · 8 THREADS · 3 BINS



PER-THREAD PRIVATIZED (LOCAL) HISTOGRAMS, THEN REDUCE  
128 ITEMS · 8 THREADS · 3 BINS  
(EACH THREAD GETS 16 ITEMS)



SORT, THEN REDUCE BY KEY







PER-THREAD PRIVATIZED (LOCAL) HISTOGRAMS, THEN REDUCE

128 ITEMS · 8 THREADS · 3 BINS

(EACH THREAD GETS 16 ITEMS)



REDUCING 8 LOCAL HISTOGRAMS



FINAL THOUGHTS ON HISTOGRAM

- ATOMICS (2)
- PER-THREAD HISTOGRAMS, THEN REDUCE (1)
- SORT, THEN REDUCE BY KEY

256 THREADS, 8 BINS:

ATOMIC TECHNIQUE:

REDUCE TO 8-ELEMENT HISTOGRAM THEN ATOMICS

HOW MANY ATOMIC ADDS?





